

Pretrained Semantic Retrieval with Clustered Indexes for Academic Paper Suggestion*

1st Srirang Kalantri

Department of Information Technology
Pune Institute of Computer Technology
Pune, India
srirangkalantri0@gmail.com

2nd Utkarsh Brahmanekar

Department of Information Technology
Pune Institute of Computer Technology
Pune, India
utkarshak4735@gmail.com

3rd Shreeram Jadhav

Department of Information Technology
Pune Institute of Computer Technology
Pune, India
shreeramjadhav564@gmail.com

4th Kirti Digholkar

Department of Information Technology
Faculty of Information Technology, Pune Institute of Computer Technology
Pune, India
kydigholkar@pict.edu

Abstract—The recent rapid increase in scholarly publications presents enormous challenges to researchers attempting to locate the most relevant previous works. Paper recommender systems that analyze article content—including titles, abstracts, and keywords—can potentially reduce the efforts spent on manual searching by automatically suggesting relevant literature. This paper presents a snapshot of recent unsupervised and content-based methods of recommending scholarly papers, focusing particular attention on TF-IDF baselines, topic modeling (e.g., LDA), clustering (covering recent deep clustering approaches), and embedding-based semantic retrieval methods (e.g., Word2Vec, Doc2Vec, and BERT/Sentence-BERT). We provide seven in-depth analyses on prominent works, propose a taxonomy of approaches, and indicate a feasible two-stage architecture making use of offline clustering and online similarity ranking, catering to efficient, real-time query-based recommendations. The survey points out evaluation-related issues and asks for a reproducible framework towards future experimental work.

Index Terms—paper recommendation, TF-IDF, LDA, cosine similarity, BERT, content-based recommendation

I. INTRODUCTION

The number of scholarly publications has grown exponentially in many areas, and the information overflow inhibits literature review work as well as novelty evaluation by researchers [1]. Even robust academic search engines that allow indexing and keyword searching often return a long or inaccurate suggestions that fail to reflect semantic proximity when varying terms carry the same meaning [2]. Content-based recommender systems alleviate this problem by analyzing textual material in documents to detect works that are semantically equivalent without having to refer to user ratings or history of user engagement [3]. They are especially useful in the academic setting, where user feedback is scarce and direct and where automated recommendations can facilitate better literature discovery, reduce redundant work, and foster better practices in citation.

Recent advancements in the natural language processing (NLP) [4] domain—especially transformer-based contextual

embeddings [5] and neural clustering methods [6]—have remarkably improved the ability to detect semantic similarity between documents. At the same time, the traditional vector-space approaches (TF-IDF with cosine similarity) remain powerful baselines because of their efficiency and interpretability [7]. Clustering methods (e.g., KMeans [8], MiniBatchKMeans [9], and deep clustering [6]) are important in the scalability of recommendations: by clustering the corpus into thematic groups offline, a system can restrict online searching to a limited subset of the dataset, thus attaining low latency. We focus on these content-based, unsupervised approaches and summarize recent literature with the goal of establishing a reproducible, pre-trained framework for recommendations.

II. LITERATURE SUREVY

A. Scientific Paper Recommendation System: Application of Sentence Transformers and Cosine Similarity Using arXiv Data [10]

This research suggests a recommending system that makes use of the Sentence Transformer technique along with cosine similarity to solve the issue of information overload in research [10]. The system is tasked with giving relevant suggestions to users. Using the sentence transformer technique enables capturing semantic context. Cosine similarity enables determining the degree of similarity between research papers represented as vectors. Though powerful, Sentence Transformer require higher computational cost.

B. A BERT Based Hybrid Recommendation System For Academic Collaboration [11]

This paper suggests a recommender system that is a combination of TF-IDF [7] and BERT [12] embeddings. The system is based on clusters that run with a computed cosine similarity matrix, the approach attempts to combine strengths of classical keyword-based schemes with the sophisticated semantic knowledge of BERT [11]. Uses clustering to organize information, which can speed up retrieval by reducing the

search space. Integrating TF-IDF and BERT, along with clustering and cosine similarity, leads to increased computational and resource requirements.

C. Research paper classification systems based on TF-IDF and LDA schemes [13]

Kim and Gil suggest a hybrid pipeline that make use of Latent Dirichlet Allocation (LDA) [14] to find latent topics in a corpus, then extracts topic-oriented keywords to create a narrowed TF-IDF dictionary, and finally performs K-Means clustering on TF-IDF vectors to organize research papers by topic [13]. Scalability is highlighted in the paper [13], and reports a Hadoop/MapReduce-based build to process large corpora, illustrating feasibility for offline batch classification. The bag-of-words TF-IDF representation again lacks contextual semantics and synonymy, so related papers that express these with different terms fall into different clusters.

D. Text clustering based on pre-trained models and autoencoders [15]

The research present a simultaneous deep clustering framework that encodes documents with pretrained transformers then trains an autoencoder that simultaneously learns an optimally clustering low-dimensional representation and the cluster assignments themselves [15]. Pros include that the model is denoising embeddings and fine-tuning representations specifically for isolating clusters, useful for downstream retrieval if clusters form partitions. End-to-end fine-tuning is permitted where in an iterative manner the cluster assignments guide learning of representations, sharpening distinctions for strongly related topics. Cons, in turn, include higher computational and memory burden as well as being highly dependent on hyperparameters.

E. Recommendation Algorithm Using Clustering-Based UPCSim (CB-UPCSim) [16]

The authors suggest CB-UPCSim, a clustering-enforced collaborative filtering strategy in which users/items are pre-clustered by K-Means to limit neighbor search range, then, within clusters, a profile-correlation similarity (UPCSim) [17] is used for prediction [16]. Tests on MovieLens demonstrate reduced run-time with better-than-baseline predictive accuracy. The paper's [16] key contribution is that pre-clustering entities can substantially cut online computation while maintaining accuracy, with direct applicability to content search in large paper corpora. Limitation is that it is designed with rating data and collaborative signals, not pure content features, so applying it to document content search involves adaptation.

F. Diversity Focused Content-Based Filtering Recommendation System [18]

This paper aims at diversifying recommendations in content-based filtering systems through the application of similarity threshold and clustering. The approach adopts cosine similarity in ascertaining the range of similarity in regard to user preference. Clustering is used in comparing the performance

of the model with and without the labels of the clusters. The findings present that diversifying by the application of similarity threshold is very effective in recommending better products [18]. Avoid overspecialization in recommendation by diversifying. Although the paper does not specifically discuss the tradeoff between relevance and diversification.

G. Clustering Based Recommender System for Compilation of Research Papers [19]

This paper suggests a recommend system that applies subspace clustering techniques to address scalability problems that are encountered in collaborative filtering for recommending scholarly papers. It is concerned with scanning the researcher-paper matrix in search of relationships among researchers and tries to give a list of articles that is well curated through reviews from experts who work in related disciplines [19]. The system provides a list of scholarly papers for further academic activities. Handles the scalability challenge associated with information overload and addresses difficulty faced in finding relevant information/work by researchers. Could limit recommendations for new papers or users.

A comparative overview of the surveyed approaches is summarized in TABLE I.

III. EXISTING METHODOLOGIES AND PREVALENT TECHNOLOGIES

A. Recommendation System Methodologies

a) Content-Based Recommendation System: Content-based systems of item recommendation function by analyzing item attributes in items preferred by the user and discerning their common features. These determined preferences are then stored in a user profile such that they can recommend items that contain a high similarity to that profile [20]. This system successfully identifies particular user interests and is capable of recommending rare items that would not gain much popularity. TF-IDF [7] and cosine similarity [10] are also commonly used in content-based systems to determine research paper relevance to a user query or profile. However, content-based systems require heavy domain knowledge for manual feature representation and are limited to recommending based only on prior interests, thereby inhibiting a user's preference from growing [20].

b) Collaborative Filtering-Based Recommendation Systems: Collaborative Filtering (CF) methods utilize the behaviors exhibited by a collective of users to generate recommendations for other individuals [20]. Such systems are frequently implemented within large-scale data processing frameworks due to their inherent characteristics that enable parallel processing [20]. The primary classifications of CF are user-based CF, which offers recommendations predicated on products favored by similar users, and item-based CF, which forecasts user ratings for items based on the similarities among those items. Generally, the later produces superior results compared to former CF [20], as the latter often experiences challenges related to sparsity and scalability. Nevertheless,

TABLE I
COMPARATIVE SUMMARY OF RELATED WORKS

Ref.	Core Technique(s)	Primary Goal / Contribution	Stated Limitation / Con
[10]	Sentence Transformers, Cosine Similarity	Provide semantically relevant suggestions to solve information overload.	High computational cost.
[11]	Hybrid: TF-IDF + BERT, Clustering	Bridge keyword-based methods with semantic knowledge from BERT.	Increased computational and resource requirements.
[13]	Hybrid: LDA + TF-IDF, K-Means	Batch classification for large corpora using a scalable pipeline.	Lacks contextual semantics (bag-of-words issue).
[15]	Pre-trained Transformers, Autoencoders	Learn document representation and clustering jointly for deep clustering.	Computationally intensive and dependent on hyperparameters.
[16]	K-Means Clustering, Collaborative Filtering (UPCSim)	Reduce computation time in collaborative filtering by pre-clustering.	Works only on user-item rating data, not on content.
[18]	Content-Based Filtering, Clustering, Cosine Similarity	Improve recommendation diversity through clustering-based filtering.	Does not discuss the trade-off between diversity and relevance.
[19]	Subspace Clustering, Collaborative Filtering	Address scalability challenges in collaborative filtering for academic data.	Limited generalization for new users or papers (cold start).

Note - CF = Collaborative Filtering; TF-IDF = Term Frequency-Inverse Document Frequency; LDA = Latent Dirichlet Allocation; BERT = Bidirectional Encoder Representations from Transformers; UPCSIm = User Profile Clustering Similarity.

both methodologies can face cold-start issues concerning users and items lacking adequate interaction data [20].

As shown in Fig. 1, the distinction between content-based and collaborative filtering-based recommendation systems is visually illustrated.

c) Knowledge-Based Recommendation Systems: This kind of systems make suggestions based on the necessary information of users, items, and their relationships [20]. It functions autonomously from user rating or purchase history, thus avoiding cold-start problems. It is particularly suitable for complex industries in which items are not commonly bought, for example, automobiles or properties [20]. However, accessing the necessary domain knowledge can pose a significant drawback for this technique [20].

d) Hybrid-Based Recommendation Systems: The hybrid recommender systems combine the strengths of various recommender methods in an effort to tackle the problems in traditional systems [20]. The systems utilize various methods of hybridization that include weighted, mixed, switching, feature combination and augmentation, cascade, and meta-level methods [20]. The most typical practice is that of combining methods of CF with content-based or knowledge-based methods in an effort to overcome sparsity, scalability, and cold-start problems [20].

A comparison of popular recommendation methodologies is summarized in TABLE II.

TABLE II
COMPARISON OF RECOMMENDATION METHODOLOGIES

Methodology	Core Mechanism	Key Strengths	Key Weaknesses
Content-Based	Analyzes attributes of items preferred by a user.	Recommends niche items; identifies specific interests.	Requires domain knowledge; causes overspecialization.
Collaborative Filtering	Utilizes collective behavior of multiple users.	Finds novel items; needs no domain knowledge.	Suffers from cold-start issues; faces data sparsity.
Knowledge-Based	Employs explicit domain knowledge and rules.	No cold-start problem; ideal for rare or high-value items.	Knowledge acquisition is time-consuming.
Hybrid-Based	Combines multiple recommendation strategies.	Balances strengths of individual approaches.	Increases implementation complexity.




Fig. 1. Illustration of (a) content-based recommendation system and (b) collaborative filtering-based recommendation system (reprinted from [20], under the Creative Commons CC BY 4.0 license).

B. Prevalent Technologies and Frameworks

a) *TF-IDF*: It is extensively utilized in content-based systems for characterizing documents as vectors, with the significance of each term signifying its relevance within a document with regard to the entire corpus [7].

b) *Embeddings*: Embedding techniques change high-dimensional discrete features such as item and user identifiers into low-dimensional continuous vectors, enhancing in turn the effectiveness of recommendation through capturing fine-grained relationships between entities [21]. Commonly used embedding methods include collaborative filtering, self-supervised learning, and graph-based methods. As an example, Word2Vec, Doc2Vec, BERT, and SBERT build paper and user profile embeddings for research papers as well as user profiles [22].

c) *Cosine Similarity*: The cosine similarity is also a good measure of the similarity degree of two vectors, for example, of document or user profile vectors. It is also utilized frequently in content-based systems as well as in collaborative filtering systems for sorting the results [10]. Fig. 2 shows conceptual workflow of cosine similarity.

d) *Clustering*: Clustering techniques are used in order to classify similar items or users, thus improving retrieval process by reducing the search space. In research paper recommendations, clustering is used for organizing academic papers, so that a faster discernment of relevant documents is achieved [22]. Fig. 3 shows compact clustering workflow.

e) *Deep Learning Architectures*: Advanced models that include recurrent neural networks and transformers, also come into the scope due to their capability of encoding and utilizing contextual information in the recommender systems [11]. Graph Neural Networks are also used to aggregate information from close papers and distinguish user preferences [22].




Fig. 2. Conceptual flow of cosine similarity




Fig. 3. Clustering workflow

IV. PROPOSED SOLUTION

The approach we take is one of cluster-aware, dual-formulation recommendation pipelines that take advantage of the efficiency of classical vector-space representations combined with the deep semantics offered by pre-trained representations. Unlike past methods depending exclusively on TF-IDF and K-Nearest Neighbors (KNN) towards quantifying paper similarity, the system takes up an off-line clustering-based online semantic re-ranking method.

A. Offline Phase: Corpus Preparation and Clustering

In the offline process, the body of research papers is preprocessed first by cleaning, tokenization, and lemmatization. Each abstract is represented by the TF-IDF scheme, and optionally by domain-specific embedding like SciBERT [23] or SentenceBERT [24]. These then are clustered by the algorithms like K-Means or MiniBatch K-Means, constituting the semantically cohesive sets of papers. Centroid vectors of each set are maintained as access points for fast similarity search.

B. Online Phase: Query Handling and Semantic Retrieval

In the online phase, when an end-user provides a paper abstract, the system retrieves the closest cluster centroid and does an in-cluster search by cosine similarity or by an approximate nearest-neighbor search by way of a library such as FAISS [25]. The papers so retrieved are re-ranked by sentence embedding so as to favor semantically correct results. This keeps the latency down compared to full-corpus search but keeps high quality of retrievals.

C. Explainability and User Interaction

A notable characteristic that sets aside the model presented is the explainability module that presents the most effective TF-IDF keywords that impacted each recommendation. The module is conducive to user confidence and transparency, something that is not provided for most deep learning-based models. The architecture also accommodates auxiliary functionalities, including the users' ability to select the number of clusters and export the recommendation outcomes.

Fig. 4 shows the proposed system architecture.




Fig. 4. System architecture.

D. Summary

By adopting the mixture of effective clustering, re-ranking by semantics, and interpretability features, the proposed approach enjoys an optimally tuned balance among scalability, accuracy, and usability. The approach surmounts the traditional TF-IDF + KNN paradigms while still enjoying feasibility toward eventual deployments on realistic data as well as the potential investigations toward maturity as an all-around

recommendation system.

REFERENCES

- [1] D. Roy and M. Dutta, "A systematic review and research perspective on recommender systems," *Journal of Big Data*, vol. 9, p. 59, 2022, doi:10.1186/s40537-022-00592-5.
- [2] J. Beel, B. Gipp, S. Langer, and C. Breiteringer, "Research-paper recommender systems: A literature survey," *International Journal on Digital Libraries*, pp. 1–34, 2015, doi:10.1007/s00799-015-0156-0.

- [3] C. C. Aggarwal, "Content-Based Recommender Systems," in *Recommender Systems: The Textbook*, Cham, Switzerland: Springer, 2016, pp. 139–166, doi:10.1007/978-3-319-29659-3_4.
- [4] B. D. Shivahare, A. K. Singh, N. Uppal, A. Rizwan, V. S. Vaathsav, and S. Suman, "Survey Paper: Study of Natural Language Processing and Its Recent Applications," in *2022 2nd International Conference on Innovative Sustainable Computational Technologies (CISCT)*, Dehradun, India, 2022, pp. 1–5, doi:10.1109/CISCT55310.2022.10046440.
- [5] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention Is All You Need," in *Proceedings of the 31st International Conference on Neural Information Processing Systems (NIPS)*, Long Beach, CA, USA, 2017, pp. 6000–6010.
- [6] J. Xie, R. Girshick, and A. Farhadi, "Unsupervised deep embedding for clustering analysis," in *Proceedings of the 33rd International Conference on Machine Learning (ICML)*, M. F. Balcan and K. Q. Weinberger, Eds., New York, NY, USA, Jun. 20–22, 2016, vol. 48, pp. 478–487. [Online]. Available: <https://proceedings.mlr.press/v48/xieb16.html>
- [7] G. Salton, A. Wong, and C. S. Yang, "A vector space model for automatic indexing," *Communications of the ACM*, vol. 18, no. 11, pp. 613–620, Nov. 1975, doi:10.1145/361219.361220.
- [8] J. MacQueen, "Some methods for classification and analysis of multivariate observations," in *Proceedings of the 5th Berkeley Symposium on Mathematical Statistics and Probability*, 1967. [Online]. Available: <https://api.semanticscholar.org/CorpusID:6278891>
- [9] D. Sculley, "Web-scale k-means clustering," in *Proceedings of the 19th International Conference on World Wide Web (WWW '10)*, Raleigh, NC, USA, 2010, pp. 1177–1178, doi:10.1145/1772690.1772862.
- [10] A. Putra, D. Putri, and A. A. K. A. Wiranatha, "Scientific paper recommendation system: Application of sentence transformers and cosine similarity using arXiv data," *Journal of Applied Informatics and Computing*, vol. 9, pp. 1374–1382, 2025, doi:10.30871/jaic.v9i4.9766.
- [11] S. Nagarajan, H. Thangaraj, V. Vashisht, E. Joshi, K. Verma, and D. Katariya, "A BERT-based hybrid recommendation system for academic collaboration," in *Proceedings of the International Conference on Artificial Intelligence and Applications*, Singapore: Springer, 2025, doi:10.1007/978-981-96-4273-1_9.
- [12] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of deep bidirectional transformers for language understanding," *arXiv preprint arXiv:1810.04805*, 2018, doi:10.48550/arXiv.1810.04805.
- [13] S.-W. Kim and J.-M. Gil, "Research paper classification systems based on TF-IDF and LDA schemes," *Human-centric Computing and Information Sciences*, vol. 9, no. 30, pp. 1–12, 2019, doi:10.1186/s13673-019-0192-7.
- [14] Z. Tong and H. Zhang, "A text mining research based on LDA topic modelling," *Computer Science & Information Technology*, vol. 6, pp. 201–210, 2016, doi:10.5121/csit.2016.60616.
- [15] Q. Xu, H. Gu, and S. Ji, "Text clustering based on pre-trained models and autoencoders," *Frontiers in Computational Neuroscience*, vol. 17, 2024. [Online]. Available: <https://api.semanticscholar.org/CorpusID:266796200>.
- [16] T. Widiyaningtyas, I. Hidayah, and T. B. Adji, "Recommendation algorithm using clustering-based UPCSIm (CB-UPCSIm)," *Computers*, vol. 10, no. 10, p. 123, 2021, doi:10.3390/computers10100123.
- [17] T. Widiyaningtyas, I. Hidayah, and T. B. Adji, "User profile correlation-based similarity (UPCSIm) algorithm in movie recommendation system," *Journal of Big Data*, vol. 8, p. 52, 2021, doi:10.1186/s40537-021-00425-x.
- [18] C. Pardede, P. R. Togatorop, A. Aritonang, and R. Lumbantoruan, "Diversity focused content-based filtering recommendation system," in *2024 2nd International Conference on Technology Innovation and Its Applications (ICTIIA)*, 2024, pp. 1–6. [Online]. Available: <https://api.semanticscholar.org/CorpusID:274450157>.
- [19] V. Verma, V. Saxena, V. Verma, and K. V. Singh, "Clustering based recommender system for compilation of research papers," *Journal of Advances in Mathematics and Computer Science*, 2024. [Online]. Available: <https://api.semanticscholar.org/CorpusID:271759654>.
- [20] Z. Xia, A. Sun, J. Xu, Y. Peng, R. Ma, and M. Cheng, "Contemporary recommendation systems on big data and their applications: A survey," *IEEE Access*, vol. 12, pp. 196914–196928, 2024, doi:10.1109/ACCESS.2024.3517492.
- [21] S. Li, H. Guo, X. Tang, R. Tang, L. Hou, R. Li, and R. Zhang, "Embedding compression in recommender systems: A survey," *ACM Computing Surveys*, vol. 56, no. 5, art. no. 130, pp. 1–21, May 2024, doi:10.1145/3637841.
- [22] C. K. Kreutz and R. Schenkel, "Scientific paper recommendation systems: A literature review of recent publications," *International Journal on Digital Libraries*, vol. 23, pp. 335–369, 2022, doi:10.1007/s00799-022-00339-w.
- [23] I. Beltagy, K. Lo, and A. Cohan, "SciBERT: A pretrained language model for scientific text," in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, Hong Kong, China, Nov. 2019, pp. 3615–3620, doi:10.18653/v1/D19-1371.
- [24] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-networks," in *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019. [Online]. Available: <https://api.semanticscholar.org/CorpusID:201646309>.
- [25] J. Johnson, M. Douze, and H. Jegou, "Billion-scale similarity search with GPUs," *IEEE Transactions on Big Data*, vol. 7, no. 3, pp. 535–547, Jul. 2021, doi:10.1109/TBDATA.2019.2921572.